

Chapter 10: Demo BoostRegressor

```
In [1]: # from google.colab import drive
# drive.mount("/content/gdrive", force_remount=True)
```

```
In [2]: # %cd '/content/gdrive/My Drive/LDS6_MachineLearning/practice/Chapter10_Boosting,
```

```
In [3]: # Load libraries
from sklearn.ensemble import AdaBoostRegressor
import pandas as pd
```

```
In [4]: # Load data
iris = pd.read_excel("Iris.xls")
iris.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
sepal.length    150 non-null float64
sepal.width     150 non-null float64
petal.length    150 non-null float64
petal.width     150 non-null float64
iris            150 non-null object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
In [5]: # 5 first samples
X = iris[["sepal.length", "sepal.width", "petal.width"]]
X.head(3)
```

Out[5]:

	sepal.length	sepal.width	petal.width
0	5.1	3.5	0.2
1	4.9	3.0	0.2
2	4.7	3.2	0.2

```
In [6]: # 5 first result
y = iris[["petal.length"]]
y.head(3)
```

Out[6]:

	petal.length
0	1.4
1	1.4
2	1.3

AdaBoostRegressor

```
In [7]: # Create adaboost-decision tree classifier object
# n_estimators: It controls the number of weak learners.
# learning_rate: Controls the contribution of weak learners in the final combination.
# There is a trade-off between learning_rate and n_estimators.
# base_estimators: It helps to specify different ML algorithm.
from sklearn.tree import DecisionTreeRegressor
ml = DecisionTreeRegressor(max_depth=4)
clf = AdaBoostRegressor(n_estimators=50,
                        base_estimator=ml,
                        learning_rate=1)
```

```
In [8]: # Train model
clf.fit(X, y)
```

c:\program files\python36\lib\site-packages\sklearn\utils\validation.py:724: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
 y = column_or_1d(y, warn=True)

```
Out[8]: AdaBoostRegressor(base_estimator=DecisionTreeRegressor(criterion='mse',
                                                                max_depth=4,
                                                                max_features=None,
                                                                max_leaf_nodes=None,
                                                                min_impurity_decrease=0.0,
                                                                min_impurity_split=None,
                                                                min_samples_leaf=1,
                                                                min_samples_split=2,
                                                                min_weight_fraction_leaf
                                                                =0.0,
                                                                presort=False,
                                                                random_state=None,
                                                                splitter='best'),
                          learning_rate=1, loss='linear', n_estimators=50,
                          random_state=None)
```

```
In [9]: X_test = [[6.4, 3.2, 1.5], [5.9, 3. , 1.8]]
y_pred = clf.predict(X_test)
y_pred
```

```
Out[9]: array([4.68235294, 4.95      ])
```

```
In [10]: # với DecisionTreeRegressor mà không có AdaBoost
clf_1 = ml.fit(X,y)
```

```
In [11]: y_pred = clf_1.predict(X_test)
y_pred
```


```
Out[11]: array([4.74375 , 5.05217391])
```

```
In [12]: X_test = pd.DataFrame(X_test)
X_test_width = X_test[2]
```

```
In [13]: import matplotlib.pyplot as plt
```

```
In [14]: # Plot the results
y_1 = clf.predict(X) # AdaBoost & DecisionTreeRegressor

plt.figure()
plt.scatter(X["petalwidth"], y, c="blue", label="training samples")
plt.scatter(X["petalwidth"], y_1, c="red",
            label="AdaBoost & DecisionTreeRegressor")
plt.scatter(X_test_width, y_pred, c="green", label="Prediction")
plt.xlabel("petalwidth")
plt.ylabel("pentallength")
plt.title("AdaBoost & DecisionTreeRegressor")
plt.legend()
plt.show()
```



XGBRegressor

```
In [15]: import xgboost as xgb
from sklearn.metrics import mean_squared_error
```

```
In [16]: xgb_model = xgb.XGBRegressor(random_state=42)
xgb_model.fit(X, y)
```

[15:20:19] WARNING: C:/Jenkins/workspace/xgboost-win64_release_0.90/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

```
Out[16]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                      colsample_bynode=1, colsample_bytree=1, gamma=0,
                      importance_type='gain', learning_rate=0.1, max_delta_step=0,
                      max_depth=3, min_child_weight=1, missing=None, n_estimators=100,
                      n_jobs=1, nthread=None, objective='reg:linear', random_state=42,
                      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
                      silent=None, subsample=1, verbosity=1)
```

```
In [17]: y_pred = xgb_model.predict(X)
print("Model r^2 score:", xgb_model.score(X,y))
```

Model r^2 score: 0.9917113219896987

```
In [18]: mse=mean_squared_error(y, y_pred)
print(mse)
```

0.025632114175041725

```
In [19]: # predict new sample
X_new = pd.DataFrame({'sepalength': [6.4,5.9],
                      'sepalwidth':[3.2,3],
                      'petalwidth':[1.5,1.8]
                      })
xgb_model.predict(X_new)
```

Out[19]: array([4.634132, 4.925619], dtype=float32)

```
In [ ]:
```